



FRIEDRICH-ALEXANDER-  
UNIVERSITÄT  
ERLANGEN-NÜRNBERG  
SCHOOL OF ENGINEERING

# Common Practices

**A. Maier**, V. Christlein, K. Breininger, Z. Yang, L. Rist, M. Nau, S. Jaganathan, C. Liu, N. Maul, L. Folle,  
K. Packhäuser, M. Zinnen

Pattern Recognition Lab, Friedrich-Alexander-Universität Erlangen-Nürnberg

April 24, 2023



# **Outline**

**Recap**

**Training Strategies**

**Optimization and Learning Rate**

**Architecture Selection and Hyperparameter Optimization**

**Ensembling**

**Class Imbalance**

**Evaluation**



FRIEDRICH-ALEXANDER-  
UNIVERSITÄT  
ERLANGEN-NÜRNBERG  
SCHOOL OF ENGINEERING

# Recap



# Training a Neural Network

- So far: all the nuts and bolts about how to train a network:
  - Fully connected and convolutional layers
  - Activation function
  - Loss function
  - Optimization
  - Regularization
- Today: Common practices on how to **choose an architecture**, **train** and **evaluate** a deep neural network.

## First Things First: Test Data



*"Ideally, the test set should be kept in a vault, and be brought out only at the end of the data analysis."*

T. Hastie, R. Tibshirani, J. Friedman: The Elements of Statistical Learning

## First Things First: Test Data (cont.)

- Overfitting is extremely easy with neural networks (see e.g. ImageNet with random labels [5]).
- True test set error/generalization error can be underestimated **substantially** when using the test set for model selection!
- Attention: Choosing the architecture is the first element in model selection  
→ should never be done on the test set!
- Do initial experimentation on smaller subset of the dataset!



FRIEDRICH-ALEXANDER-  
UNIVERSITÄT  
ERLANGEN-NÜRNBERG  
SCHOOL OF ENGINEERING


# Training Strategies



## Before Training: Gradient Checks

Own loss function, own layer implementation etc.: Check correct computation of gradient by comparing analytic and numerical gradient.

- Use centered differences for numeric gradient.




Source: <https://fhsperformance.wordpress.com/2016/09/02/ramp-up-your-warm-ups/>



## Before Training: Gradient Checks

Own loss function, own layer implementation etc.: Check correct computation of gradient by comparing analytic and numerical gradient.

- Use centered differences for numeric gradient.
- Use relative error instead of absolute differences.




Source: <https://fhsperformance.wordpress.com/2016/09/02/ramp-up-your-warm-ups/>

## Before Training: Gradient Checks

Own loss function, own layer implementation etc.: Check correct computation of gradient by comparing analytic and numerical gradient.

- Use centered differences for numeric gradient.
- Use relative error instead of absolute differences.
- Numerics:
  - Use double precision for checking.
  - Temporarily scale loss function if you observe very small values ( $< 1e - 9$ ).
  - Choose  $h$  appropriately.



Source: <https://fhsperformance.wordpress.com/2016/09/02/ramp-up-your-warm-ups/>


## Before Training: Gradient Checks (cont.)

Additional recommendations:

- Use only a few datapoints → less issues with non-differentiable parts of the loss function.
- Train the network for a short period of time before performing gradient checks.
- Check gradient first without, then with regularization terms.
- Turn off data augmentation and dropout.

## Before Training: Check Initialization and Loss


- Goal: Check correct random initialization of layers.
- Compute the loss for each class on the **untrained network**, with regularization turned off.



Source: <https://commons.wikimedia.org/>

## Before Training: Check Initialization and Loss


- Goal: Check correct random initialization of layers.
- Compute the loss for each class on the **untrained network**, with regularization turned off.
- Compare loss with loss achieved when deciding for a class randomly (**chance**).



Source: <https://commons.wikimedia.org/>

## Before Training: Check Initialization and Loss

- Goal: Check correct random initialization of layers.
- Compute the loss for each class on the **untrained network**, with regularization turned off.
- Compare loss with loss achieved when deciding for a class randomly (**chance**).
- Repeat with multiple random initializations.



Source: <https://commons.wikimedia.org/>

## Before Training: Training!

- Goal: Check whether the architecture is **in general** capable to learn the task.
- Before training the network on the full training data set, take a small subset (5-20 samples) and try to **overfit** the network to get zero loss.
- Optionally: Turn off regularization that may hinder overfitting.

## Before Training: Training!

- Goal: Check whether the architecture is **in general** capable to learn the task.
- Before training the network on the full training data set, take a small subset (5-20 samples) and try to **overfit** the network to get zero loss.
- Optionally: Turn off regularization that may hinder overfitting.
- If the network cannot overfit:
  - Bug in the implementation.
  - Model too small → increase number of parameters.
  - Model not suitable for the task.




## Before Training: Training!

- Goal: Check whether the architecture is **in general** capable to learn the task.
- Before training the network on the full training data set, take a small subset (5-20 samples) and try to **overfit** the network to get zero loss.
- Optionally: Turn off regularization that may hinder overfitting.
- If the network cannot overfit:
  - Bug in the implementation.
  - Model too small → increase number of parameters.
  - Model not suitable for the task.
- Also: Get a first idea about how the data, loss and network behave.

## During Training: Monitor loss function


- Recap:



- Check learning rate (→more in a bit).
- Identify large jumps in the learning curve.


## During Training: Monitor loss function

- Recap:



- Check learning rate (→more in a bit).
- Identify large jumps in the learning curve.
- Very noisy curves →increase batch size.


## During Training: Monitor Validation Loss



- Monitor amount of overfitting of the network.
- If training and validation loss diverge: overfitting → increase regularization/ early stopping
- If training and validation loss are close but high: underfitting → decrease regularization/ increase model size
- Save intermediate models if you want to use them for testing!

## During Training: Monitor Weights and Activations

- Track relative magnitude of the weight update: Should be in a sensible range (approx.  $1e-3$ ).
- Convolutional layers: check filters of the first few layers. Should develop towards smooth and regular filters.
- Check for very large or saturated activations ( $\rightarrow$  dying ReLUs)



Source: <https://cs231n.github.io/neural-networks-3/>



FRIEDRICH-ALEXANDER-  
UNIVERSITÄT  
ERLANGEN-NÜRNBERG  
SCHOOL OF ENGINEERING

# Optimization and Learning Rate




## Choosing an Optimizer

- Batch gradient descent: Requires large memory, too slow, too few updates.
- Stochastic gradient descent (SGD): loss function and gradient become very noisy if only one/few samples are used.
- SGD with mini-batches: “best of both worlds”
  - Frequent, more stable updates.
  - Gradient noisy enough to escape local minima.
  - Adapting mini-batch size yields smoother/more noisy gradient.
- Addition of momentum prevents oscillations and speeds up optimization.
- Effect of hyper-parameters relatively straight forward.
- Recommendation: Start with Mini-Batch SGD + momentum.
- For faster convergence speed → ADAM.

## Learning rate: Observing the loss curve

- Learning rate  $\eta$  has a large impact on the successful training of a network.
- For almost all gradient based optimizers,  $\eta$  has to be set.
- Effect of learning rate is often directly observable in the loss curve.



- But this is a very simplified view!
- We want an adaptive learning rate: Progressively smaller steps to find the optimum
- **Annealing** the learning rate.



## Annealing the Learning Rate

- In deep learning context often known as **learning rate decay**.
- Decay means yet another hyper-parameter.
- Need avoid oscillation as well as a too fast cool down!
- Decay strategies:
  - **Stepwise decay**: Every  $n$  epochs, reduce learning rate by a certain factor, e.g. 0.5, or by a constant value, e.g. 0.01.  
Variant: Reduce learning rate when validation error stagnates.
  - **Exponential decay**: At epoch  $t$ :  $\eta = \eta_0 e^{-kt}$  with  $k$  controlling the decay.
  - **1/t-decay**: At epoch  $t$ :  $\eta = \eta_0 / (1 + kt)$ .
- Stepwise decay most common: hyper-parameters are easy to interpret.
- Second-order methods are currently uncommon in practice, as they do not scale as well.

**NEXT TIME**

**ON DEEP LEARNING**



FRIEDRICH-ALEXANDER-  
UNIVERSITÄT  
ERLANGEN-NÜRNBERG  
SCHOOL OF ENGINEERING

# Common Practices - Part 2

**A. Maier**, V. Christlein, K. Breininger, Z. Yang, L. Rist, M. Nau, S. Jaganathan, C. Liu, N. Maul, L. Folle,  
K. Packhäuser, M. Zinnen

Pattern Recognition Lab, Friedrich-Alexander-Universität Erlangen-Nürnberg

April 24, 2023





FRIEDRICH-ALEXANDER-  
UNIVERSITÄT  
ERLANGEN-NÜRNBERG  
SCHOOL OF ENGINEERING

# Architecture Selection and Hyperparameter Optimization



## Reminder



**Test data → vault!**

# Hyperparameter optimization

Neural networks have an enormous amount of hyperparameters.

- Architecture:
  - Number of layers & number of nodes per layer
  - Activation function
  - ...
- Optimization
  - Initialization
  - Loss function
  - Optimizer (SGD, Momentum, ADAM, ...)
  - Learning rate, decay & batch size
  - ...
- Regularization
  - Regularizer, e.g.,  $L_2$ -,  $L_1$ -loss
  - Batch Normalization?
  - Dropout?
  - ...
- ...

## Choosing Architecture and Loss Function

- First step: Think about the problem and the data:
  - How could the features look like?
  - What kind of spatial correlation do you expect?
  - What data augmentation makes sense?
  - How will the classes be distributed?
  - What is important regarding the target application?
- Start with simple architectures and loss functions.
- Do your research: Try **well-known** models first and foremost!
- If you change/adapt the architecture: Find reasons why the network should perform better.

## Hyperparameter search

- Learning rate, decay, regularization/dropout etc. can be tuned more easily.
- Still, networks can take days/weeks to train!

Source: [2]



## Hyperparameter search

- Learning rate, decay, regularization/dropout etc. can be tuned more easily.
- Still, networks can take days/weeks to train!
- Search for hyperparameters using a log scale (e.g.,  $\eta \in \{0.1, 0.01, 0.001\}$ ).

Source: [2]


## Hyperparameter search

- Learning rate, decay, regularization/dropout etc. can be tuned more easily.
- Still, networks can take days/weeks to train!
- Search for hyperparameters using a log scale (e.g.,  $\eta \in \{0.1, 0.01, 0.001\}$ ).
- Options: Grid search or random search:

Source: [2]

## Hyperparameter search

- Learning rate, decay, regularization/dropout etc. can be tuned more easily.
- Still, networks can take days/weeks to train!
- Search for hyperparameters using a log scale (e.g.,  $\eta \in \{0.1, 0.01, 0.001\}$ ).
- Options: Grid search or random search:
  - Use random search instead of grid search [2]:
    - Easier to implement.
    - Better exploration of parameters that have strong influence on the result.



Source: [2]

## Hyperparameter search: Coarse to fine search

- Hyperparameters highly interdependent.
- Optimize on a coarse to fine scale:
  - Training network only for a few epochs.
  - Bring all hyperparameters in sensible ranges.
  - Then refine using random/grid-search.



FRIEDRICH-ALEXANDER-  
UNIVERSITÄT  
ERLANGEN-NÜRNBERG  
SCHOOL OF ENGINEERING

# Ensembling



## Concept

- So far we have always considered a **single** classifier. Can't we get better by using **many**?
- Assume  $N$  classifiers **independently** performing a correct prediction with probability  $1 - p$
- The probability of seeing  $k$  errors is:


$$\binom{N}{k} p^k (1 - p)^{N-k},$$

known as binomial distribution


- So the probability of a majority  $k > \frac{N}{2}$  to be wrong is:

$$\sum_{k > \frac{N}{2}}^N \binom{N}{k} p^k (1 - p)^{N-k}$$

## Binomial distribution for increasing $N$




## Binomial distribution for increasing $N$



- $\sum_{k > \frac{N}{2}} \binom{N}{k} p^k (1-p)^{N-k}$   
**monotonically**  
**decreasing for  $N \rightarrow \infty$**




## Binomial distribution for increasing $N$



- $\sum_{k > \frac{N}{2}} \binom{N}{k} p^k (1-p)^{N-k}$   
**monotonically**  
**decreasing for  $N \rightarrow \infty$**
- **Accuracy  $\rightarrow 1!$**

## Binomial distribution for increasing $N$



- $\sum_{k > \frac{N}{2}} \binom{N}{k} p^k (1 - p)^{N-k}$   
**monotonically decreasing** for  $N \rightarrow \infty$
- **Accuracy**  $\rightarrow 1!$
- The big assumption here is **independence**

## Concept (cont.)

### Ensembling

- Produce  $N$  **independent** classifiers/regressors
- **Combine** their **predictions** by majority/averaging

### How to produce the components?

## Concept (cont.)


### Ensembling

- Produce  $N$  **independent** classifiers/regressors
- **Combine** their **predictions** by majority/averaging

### How to produce the components?


- Different **models**

## Local Minima




- Can we use multiple local minima we get during training?

## Local Minima



- Can we use multiple local minima we get during training?
- Combine models across optimization process

## Local Minima



- Can we use multiple local minima we get during training?
- Combine models across optimization process
- Can be combined with a **cyclic** learning rate

## Concept (cont.)

### Ensembling

- Produce  $N$  **independent** classifiers/regressors
- **Combine** their **predictions** by majority/averaging

### How to produce the components?

- Different **models**
- Different model **checkpoints**



## Concept (cont.)

### Ensembling

- Produce  $N$  **independent** classifiers/regressors
- **Combine** their **predictions** by majority/averaging

### How to produce the components?

- Different **models**
- Different model **checkpoints**
- **Moving average** of  $w$  [6]

## Concept (cont.)

### Ensembling

- Produce  $N$  **independent** classifiers/regressors
- **Combine** their **predictions** by majority/averaging

### How to produce the components?

- Different **models**
- Different model **checkpoints**
- **Moving average** of  $w$  [6]
- Different **methods**

## Concept (cont.)

### Ensembling

- Produce  $N$  **independent** classifiers/regressors
- **Combine** their **predictions** by majority/averaging

### How to produce the components?

- Different **models**
- Different model **checkpoints**
- **Moving average** of  $w$  [6]
- Different **methods**

→ **Easy performance boost if you need just a bit more**

**NEXT TIME**

**ON DEEP LEARNING**



FRIEDRICH-ALEXANDER-  
UNIVERSITÄT  
ERLANGEN-NÜRNBERG  
SCHOOL OF ENGINEERING

# Common Practices - Part 3

**A. Maier**, V. Christlein, K. Breininger, Z. Yang, L. Rist, M. Nau, S. Jaganathan, C. Liu, N. Maul, L. Folle,  
K. Packhäuser, M. Zinnen

Pattern Recognition Lab, Friedrich-Alexander-Universität Erlangen-Nürnberg

April 24, 2023





FRIEDRICH-ALEXANDER-  
UNIVERSITÄT  
ERLANGEN-NÜRNBERG  
SCHOOL OF ENGINEERING

# Class Imbalance



## Motivation

- Often, different classes are available with very different frequencies in the data set.
- Big challenge for machine learning algorithms.



## Motivation

- Often, different classes are available with very different frequencies in the data set.
- Big challenge for machine learning algorithms.
- Example 1: Fraud detection
  - Out of 10000 transactions, 9999 are genuine and 1 is fraudulent:





## Motivation

- Often, different classes are available with very different frequencies in the data set.
- Big challenge for machine learning algorithms.
- Example 1: Fraud detection
  - Out of 10000 transactions, 9999 are genuine and 1 is fraudulent:
    - Classifying every transaction as genuine:  
99.99% accuracy




## Motivation

- Often, different classes are available with very different frequencies in the data set.
- Big challenge for machine learning algorithms.
- Example 1: Fraud detection
  - Out of 10000 transactions, 9999 are genuine and 1 is fraudulent:
    - Classifying every transaction as genuine: 99.99% accuracy
    - Misclassifying 1 out of 100 genuine transactions: 99% accuracy




## Motivation (cont.)

- Task: Detect mitotic cells for tumor diagnostics [1].
- Problem: Mitotic cells only make up a very small portion of cells in tissues.




## Motivation (cont.)

- Task: Detect mitotic cells for tumor diagnostics [1].
- Problem: Mitotic cells only make up a very small portion of cells in tissues.
- Data of a certain class is seen much less during training.




## Motivation (cont.)

- Task: Detect mitotic cells for tumor diagnostics [1].
- Problem: Mitotic cells only make up a very small portion of cells in tissues.
- Data of a certain class is seen much less during training.
- Measures like accuracy,  $L_2$  norm, cross-entropy do not show imbalance.




## Resampling Strategies for Class Imbalance

- Idea: Balance class frequencies by sampling classes differently.




## Resampling Strategies for Class Imbalance

- Idea: Balance class frequencies by sampling classes differently.
- **Undersampling:**
  - In each iteration, take a subset of the overrepresented class.
  - Samples of all classes are now presented to the network equally often.



## Resampling Strategies for Class Imbalance


- Idea: Balance class frequencies by sampling classes differently.
- **Undersampling:**
  - In each iteration, take a subset of the overrepresented class.
  - Samples of all classes are now presented to the network equally often.
  - Disadvantage: Not all available data is used for training and can lead to underfitting.






## Resampling Strategies for Class Imbalance (cont.)

- **Oversampling:**
  - Use sample from underrepresented class multiple times.
  - All available data can be used.




## Resampling Strategies for Class Imbalance (cont.)

- **Oversampling:**
  - Use sample from underrepresented class multiple times.
  - All available data can be used.
  - Disadvantage: Can lead to overfitting.



## Resampling Strategies for Class Imbalance (cont.)

- **Oversampling:**
  - Use sample from underrepresented class multiple times.
  - All available data can be used.
  - Disadvantage: Can lead to overfitting.
- Also possible: Combine Under- and Oversampling.



## Resampling Strategies for Class Imbalance (cont.)

- More advanced resampling strategies available that try to avoid the shortcomings of simple under-/oversampling, e.g., Synthetic Minority Over-Sampling Technique (SMOTE).
- Rather uncommon in deep learning.

## Resampling Strategies for Class Imbalance (cont.)

- More advanced resampling strategies available that try to avoid the shortcomings of simple under-/oversampling, e.g., Synthetic Minority Over-Sampling Technique (SMOTE).
- Rather uncommon in deep learning.
- Underfitting caused undersampling can be reduced by taking a different subset after each epoch.

## Resampling Strategies for Class Imbalance (cont.)

- More advanced resampling strategies available that try to avoid the shortcomings of simple under-/oversampling, e.g., Synthetic Minority Over-Sampling Technique (SMOTE).
- Rather uncommon in deep learning.
- Underfitting caused undersampling can be reduced by taking a different subset after each epoch.
- Data augmentation can help to reduce overfitting for underrepresented class.

## Class imbalance: Adapt the Loss Function

- Instead of “fixing” the data, adapt the loss function to be stable with respect to class imbalance.

## Class imbalance: Adapt the Loss Function

- Instead of “fixing” the data, adapt the loss function to be stable with respect to class imbalance.
- Weight loss with inverse class frequency, e.g., weighted cross entropy:

$$L(\mathbf{y}, \hat{\mathbf{y}}) = -w_k \log(\hat{y}_k) |_{y_k=1} \quad (1)$$



## Class imbalance: Adapt the Loss Function

- Instead of “fixing” the data, adapt the loss function to be stable with respect to class imbalance.
- Weight loss with inverse class frequency, e.g., weighted cross entropy:

$$L(\mathbf{y}, \hat{\mathbf{y}}) = -w_k \log(\hat{y}_k)|_{y_k=1} \quad (1)$$

- More common in segmentation problems: Dice-loss based on Dice coefficient.

## Class imbalance: Adapt the Loss Function

- Instead of “fixing” the data, adapt the loss function to be stable with respect to class imbalance.
- Weight loss with inverse class frequency, e.g., weighted cross entropy:

$$L(\mathbf{y}, \hat{\mathbf{y}}) = -w_k \log(\hat{y}_k) |_{y_k=1} \quad (1)$$

- More common in segmentation problems: Dice-loss based on Dice coefficient.
- Instead of class frequency, weights can be adapted with regards to other considerations.

**NEXT TIME**

**ON DEEP LEARNING**



FRIEDRICH-ALEXANDER-  
UNIVERSITÄT  
ERLANGEN-NÜRNBERG  
SCHOOL OF ENGINEERING

# Common Practices - Part 4

**A. Maier**, V. Christlein, K. Breininger, Z. Yang, L. Rist, M. Nau, S. Jaganathan, C. Liu, N. Maul, L. Folle,  
K. Packhäuser, M. Zinnen

Pattern Recognition Lab, Friedrich-Alexander-Universität Erlangen-Nürnberg

April 24, 2023





FRIEDRICH-ALEXANDER-  
UNIVERSITÄT  
ERLANGEN-NÜRNBERG  
SCHOOL OF ENGINEERING

# Evaluation



## Performance evaluation

- Network was trained on training set, hyper-parameters estimated on the validation set.

## Performance evaluation

- Network was trained on training set, hyper-parameters estimated on the validation set.
- Evaluate generalization performance on previously unseen data: the test set.

## Performance evaluation

- Network was trained on training set, hyper-parameters estimated on the validation set.
  - Evaluate generalization performance on previously unseen data: the test set.
- We can now open the vault!



Source: [de.disney.wikia.com/wiki/Dagobert\\_Duck](https://de.disney.wikia.com/wiki/Dagobert_Duck)



## Of All Things the Measure is Man [8]

- Protagoras of Abdera (c.490 - c.420 BCE)
- Data is annotated and labeled by humans.
- During training, all labels are assumed to be correct ⚡ “to err is human”
- Additionally: Ambiguous data.

## Of All Things the Measure is Man [8]


- Protagoras of Abdera (c.490 - c.420 BCE)
- Data is annotated and labeled by humans.
- During training, all labels are assumed to be correct ⚡ “to err is human”
- Additionally: Ambiguous data.
- Multiple human voters: Take mean (if possible) or majority vote.

## Of All Things the Measure is Man [8]

- Protagoras of Abdera (c.490 - c.420 BCE)
- Data is annotated and labeled by humans.
- During training, all labels are assumed to be correct ⚡ “to err is human”
- Additionally: Ambiguous data.
- Multiple human voters: Take mean (if possible) or majority vote.
- Steidl et al. 2005: Entropy-based measure that takes “confusions” of human reference labelers into account.
  - Humans confuse certain classes with each other more (Angry vs. Happy/Angry vs. Annoyed)
  - Mistakes by the classifier are less severe if the same classes are confused by humans.

## Performance measures


- Classification problem → classification measures:



Source: <https://commons.wikimedia.org/>

## Performance measures


- Classification problem → classification measures:
- Binary classification problem:
  - True/False Positives: TP/FP
  - True/False Negatives: TN/FN



Source: <https://commons.wikimedia.org/>

## Performance measures


- Classification problem → classification measures:
- Binary classification problem:
  - True/False Positives: TP/FP
  - True/False Negatives: TN/FN
- Accuracy:  $ACC = \frac{TP+TN}{P+N}$
- Precision/pos. predictive value:  
 $precision = \frac{TP}{TP+FP}$
- Recall/true positive value:  $recall = \frac{TP}{TP+FN}$
- Specificity/true negative value:  
 $specificity = \frac{TN}{TN+FP}$
- F1-score:  $F1 = 2 \cdot \frac{TPV \cdot PPV}{TPV+PPV}$



Source: <https://commons.wikimedia.org/>

## Performance measures

- Classification problem → classification measures:
- Binary classification problem:
  - True/False Positives: TP/FP
  - True/False Negatives: TN/FN
- Accuracy:  $ACC = \frac{TP+TN}{P+N}$
- Precision/pos. predictive value:  
precision =  $\frac{TP}{TP+FP}$
- Recall/true positive value: recall =  $\frac{TP}{TP+FN}$
- Specificity/true negative value:  
specificity =  $\frac{TN}{TN+FP}$
- F1-score:  $F1 = 2 \cdot \frac{TPV \cdot PPV}{TPV + PPV}$
- Receiver operating characteristic (ROC) curve.



Source: <https://commons.wikimedia.org/>

## Performance measures: Multiclass classification

- Adapted versions of measures mentioned above.
- Top- $K$  error: True class label is not in the  $K$  classes with the highest prediction score.
- Common: Top-1 and Top-5 error.
- Example: ImageNet performance usually measured with Top-5 error.



## Cross Validation

- $k$ -fold cross validation:
  - Split data in  $k$  folds
  - Use  $k - 1$  folds as training data, test on fold  $k$
  - Repeat  $k$  times.
- Rather uncommon in deep learning due to long training times.
- Can be used for hyperparameter estimation (nested!), or to evaluate stability of (hyper-)parameters.

## Cross Validation

- $k$ -fold cross validation:
  - Split data in  $k$  folds
  - Use  $k - 1$  folds as training data, test on fold  $k$
  - Repeat  $k$  times.
- Rather uncommon in deep learning due to long training times.
- Can be used for hyperparameter estimation (nested!), or to evaluate stability of (hyper-)parameters.
- Underestimates variance of results: Training runs are not independent.

## Cross Validation

- $k$ -fold cross validation:
  - Split data in  $k$  folds
  - Use  $k - 1$  folds as training data, test on fold  $k$
  - Repeat  $k$  times.
- Rather uncommon in deep learning due to long training times.
- Can be used for hyperparameter estimation (nested!), or to evaluate stability of (hyper-)parameters.
- Underestimates variance of results: Training runs are not independent.
- Attention: almost always additional bias (architecture selection, hyperparameters).

## Cross Validation

- $k$ -fold cross validation:
    - Split data in  $k$  folds
    - Use  $k - 1$  folds as training data, test on fold  $k$
    - Repeat  $k$  times.
  - Rather uncommon in deep learning due to long training times.
  - Can be used for hyperparameter estimation (nested!), or to evaluate stability of (hyper-)parameters.
  - Underestimates variance of results: Training runs are not independent.
  - Attention: almost always additional bias (architecture selection, hyperparameters).
  - Even without cross-validation: Training is a highly stochastic process.
- Retrain network multiple times and report average performance and standard deviation.

## Comparing Classifiers

- Example: Is my new method with 91.5% accuracy better than the state-of-the-art with 90.9%?
- Training a neural network is a stochastic process.
- Simply comparing two (or more) numbers yields biased results!

## Comparing Classifiers

- Example: Is my new method with 91.5% accuracy better than the state-of-the-art with 90.9%?
- Training a neural network is a stochastic process.
- Simply comparing two (or more) numbers yields biased results!
- Actual question: Is there a **significant** difference between classifiers?

## Comparing Classifiers

- Example: Is my new method with 91.5% accuracy better than the state-of-the-art with 90.9%?
  - Training a neural network is a stochastic process.
  - Simply comparing two (or more) numbers yields biased results!
  - Actual question: Is there a **significant** difference between classifiers?
  - Run training for each method/network multiple times.
- Determine whether performance is significantly different e.g. **Student's t-test!**

## Comparing Classifiers

- Example: Is my new method with 91.5% accuracy better than the state-of-the-art with 90.9%?
  - Training a neural network is a stochastic process.
  - Simply comparing two (or more) numbers yields biased results!
  - Actual question: Is there a **significant** difference between classifiers?
  - Run training for each method/network multiple times.
- Determine whether performance is significantly different e.g. **Student's t-test!**
- Compares two normally distributed data sets with equal variance.
  - Determines whether the means are significantly different with respect to a **significance level**  $\alpha$  (e.g. 0.05 or 0.01).



## Comparing Classifiers: Bonferroni Correction

- Interpretation: The probability that this difference is caused by **chance**  $< \alpha$ .
- If we compare multiple classifiers, this chance can rise significantly due to **multiple comparisons**!

## Comparing Classifiers: Bonferroni Correction

- Interpretation: The probability that this difference is caused by **chance**  $< \alpha$ .
- If we compare multiple classifiers, this chance can rise significantly due to **multiple comparisons**!
- Correct for multiple tests using Bonferroni correction:
  - For  $n$  tests with significance level  $\alpha$ , the total risk is  $n \cdot \alpha$ .
  - To reach a total significance level of  $\alpha$ , choose adjusted  $\alpha' = \alpha/n$  for each individual test.

## Comparing Classifiers: Bonferroni Correction

- Interpretation: The probability that this difference is caused by **chance**  $< \alpha$ .
- If we compare multiple classifiers, this chance can rise significantly due to **multiple comparisons**!
- Correct for multiple tests using Bonferroni correction:
  - For  $n$  tests with significance level  $\alpha$ , the total risk is  $n \cdot \alpha$ .
    - To reach a total significance level of  $\alpha$ , choose adjusted  $\alpha' = \alpha/n$  for each individual test.
- Assumes independence between tests: Pessimistic estimation of significance.
- More accurate, but incredibly time-consuming: Permutation tests [3].

## Summary

- Check your implementation before training: Gradient, initialization, ...
- Monitor training process continuously: training/validation loss, weights, activations.
- Stick to established architectures before reinventing the wheel.
- Experiment with few data sets, keep your evaluation data safe until evaluation.
- Decay the learning rate over time.
- Do random search (not grid search) for hyperparameters.
- Perform model ensembling for better performance.
- Check for significance when comparing classifiers.

**NEXT TIME**

**ON DEEP LEARNING**

## Coming Up

Evolution of neural network architectures:

- From deep networks to deeper networks.
- From “sparse” to dense connections.
- LeNet, GoogLeNet, ResNet, ...

## Further Reading

- Link SGD Tricks by Leon Bottou.
- Link: Interesting loss functions.
- Link: Practical recommendations by Yoshua Bengio (from 2012).



FRIEDRICH-ALEXANDER-  
UNIVERSITÄT  
ERLANGEN-NÜRNBERG  
SCHOOL OF ENGINEERING

# References





## References I

- [1] M. Aubreville, M. Krappmann, C. Bertram, et al. “A Guided Spatial Transformer Network for Histology Cell Differentiation”. In: [ArXiv e-prints](#) (July 2017). arXiv: 1707.08525 [cs.CV].
- [2] James Bergstra and Yoshua Bengio. “Random Search for Hyper-parameter Optimization”. In: [J. Mach. Learn. Res.](#) 13 (Feb. 2012), pp. 281–305.
- [3] Jean Dickinson Gibbons and Subhabrata Chakraborti. “Nonparametric statistical inference”. In: [International encyclopedia of statistical science](#). Springer, 2011, pp. 977–979.
- [4] Yoshua Bengio. “Practical recommendations for gradient-based training of deep architectures”. In: [Neural networks: Tricks of the trade](#). Springer, 2012, pp. 437–478.

## References II

- [5] Chiyuan Zhang, Samy Bengio, Moritz Hardt, et al. “Understanding deep learning requires rethinking generalization”. In: [arXiv preprint arXiv:1611.03530](#) (2016).
- [6] Boris T Polyak and Anatoli B Juditsky. “Acceleration of stochastic approximation by averaging”. In: [SIAM Journal on Control and Optimization](#) 30.4 (1992), pp. 838–855.
- [7] Prajit Ramachandran, Barret Zoph, and Quoc V. Le. “Searching for Activation Functions”. In: [CoRR abs/1710.05941](#) (2017). arXiv: 1710.05941.
- [8] Stefan Steidl, Michael Levit, Anton Batliner, et al. “Of All Things the Measure is Man: Automatic Classification of Emotions and Inter-labeler Consistency”. In: [Proc. of ICASSP](#). IEEE “Institute of Electrical and Electronics Engineers, Mar. 2005.