

DL Solution SS22



I can neither confirm nor deny details of any operation without the Secretary's approval.

1 Single Choice Questions (10P)

For each of the following questions, mark the **one correct choice**. Each question has **only one** correct option. No explanation is required.

Question 1.1

1 P.

Which of the following techniques is NOT a common practice to tackle class imbalance?

- ☐ Weighting loss with inverse class frequency
- ☐ Data augmentation
- ☐ Oversampling
- ☒ Batch normalization

Question 1.2

1 P.

What is ensembling?

- ☐ A visualization technique where different areas in the input image are occluded
- ☐ A network with multiple outputs for different tasks
- ☐ Parallel paths in a network each including a different kernel size respectively to let the network decide on the best kernel parameters
- ☒ Multiple networks trained on the same task where the prediction is a majority vote

Question 1.3

1 P.

Which option can be used to create synthetic samples?

- ☐ The discriminator part of a Generative Adversarial Network
- ☐ The encoder part of a Variational Autoencoder
- ☒ The decoder part of a Variational Autoencoder
- ☐ The encoder part of a Generative Adversarial Network

Question 1.4

1 P.

Imagine that you want to generate a sequence of words with your RNN after you finished the training. Which of the following sampling strategies does not exist for that purpose?

- ☐ Greedy sampling
- ☐ Beam sampling
- ☒ Recognition sampling
- ☐ Random sampling

Question 1.5

1 P.

Which statement is false when implementing a Convolutional layer?

- ☐ The number of kernels determines the number of output channels
- ☐ Stride bigger than one reduces the image size
- ☒ A valid convolution outputs the image of the same size as its input
- ☐ The number of parameters/weights inside the kernel is independent from the image size

Question 1.6

1 P.

Why does initialization matter for the optimization of deep neural networks?

- ☒ Because the problem is often non-convex
- ☐ To avoid too many parameters
- ☐ So we always end up in the same local minimum
- ☐ Because the problem is convex

Question 1.7

1 P.

What's the main tasks of Object Detection

- ☒ Finding Bounding boxes and Classification
- ☐ Search for boundaries between different objects
- ☐ Increase the image resolution of the regions of interest
- ☐ Fill out the missing part of the image

Question 1.8

1 P.

Which statement about YOLO (You Only Look Once) and Fast R-CNN (Regional Convolutional Neural Network) is True?

- ☐ YOLO produces many object candidate suggestions that are passed to the CNN
- ☒ YOLO combines bounding box prediction and classification in one network
- ☐ Fast R-CNN is in general faster
- ☐ Fast R-CNN can do real time detection while YOLO can not

Question 1.9

1 P.

It is important to create your data set properly to avoid misclassifications during the testing. Which of the following examples is not a confound?

- ☒ Food classification task: The food images were taken with two different cameras regardless of their label
- ☐ Traffic object detection task: training images of cars were taken in daylight, while images of pedestrians were taken on rainy days
- ☐ Clothing classification task: T-shirts are worn only by females while hoodies are only worn by males
- ☐ Language classification: spanish voice samples were recorded with a different microphone than the italian samples

Question 1.10

1 P.

When handling an image classification task, you might only have a few options to tackle the problem but luckily, you have a model that was already trained on similar task. Which of the following methods could be used to make use of this trained model?

- ☒ Freeze all feature extraction layers and retrain the classification layers at the end
- ☐ Use the whole trained model and only retrain the input layer
- ☐ Only retrain the feature extraction layers and use the original classification layers
- ☐ Assess each layers and pick some layers to form a new model

2 Short Answers (9P)

Question 2.1

1 P.

Solution 2.1

cross-correlation is a convolution with flipped kernel and vice-versa.

Question 2.2

1 P.

Solution 2.2

variety of functions it can approximate / number of parameters

Question 2.3

2 P.

Solution 2.3

- ReLU is not zero-centered
- Initialization and input distribution might not be normalized
- Deeper nets \rightarrow amplified effect

Question 2.4

1 P.

Solution 2.4

It contains time-dependency / hidden states as time-dependent state between each time / it contains hidden states as "short memories" to connect between each time

Question 2.5

1 P.

Solution 2.5

it contains long term memory / It has cell state that serves as long term memory so we can model longer time series / Easier to detect long-term dependencies by using cell states

Question 2.6

1 P.

Solution 2.6

Pixel Accuracy / Mean Pixel Accuracy / Mean Intersection over Union / Frequency Weighted Intersection over Union (FWIoU)

Question 2.7

1 P.

Solution 2.7

activations generated by kernels

Question 2.8

1 P.

Solution 2.8

(each could give one point, in total one point)

- 1x1 convolutions simply calculate inner products at each position
- Simple and efficient method to decrease the size of a network
- Learns dimensionality reduction, e.g., can reduce redundancy in your feature maps

3 Feed Forward Network (9P)

Question 3.1

1 P.

Why do you use ReLU activations instead of Sigmoid activations after each fully connected layer?

Solution 3.1

because Sigmoid function saturates / ReLU does not saturate in positive area. (1)

Question 3.2

2 P.

Imagine an image size of (X, X) .

Solution 3.2

$(X^2) * Z$ (1 P.)

$(K^2) * N$ (1 P.)

Question 3.3

2 P.

Solution 3.3

Disadvantage fully connected:

- Highly correlated / not independent from each other
- Scale dependent / different size
- Intensity variations / color variation

Advantages convolutional:

- Great feature extraction
- Local connectivity - i filters
- Translational equivariance

Question 3.4

1 P.

Solution 3.4

Flatten, pooling etc. 1 P.

Question 3.5

1 P.

Solution 3.5

A simpler way to understand what the bias is: it is somehow similar to the constant b of a linear function

$$y = ax + b$$

It allows you to move the line up and down to fit the prediction with the data better.

Without b , the line always goes through the origin $(0, 0)$ and you may get a poorer fit.

Question 3.6

2 P.

Solution 3.6

- Sigmoid function (1).
- Each output entry is between 0 and 1. Softmax function will make all entries sum up to 1. (1)

4 Recurrent Networks and Backpropagation (9P)

Question 4.1

3 P.

Solution 4.1

Half point each:

$$\begin{aligned}l &= x \cdot w + b \\m &= x \cdot w + b + x \\\sigma(x) &= \frac{1}{1 + \exp(x)} \\\sigma'(x) &= \sigma(x) \cdot (1 - \sigma(x)) \\f_1(x) &= \sigma(x + wx + b) \\f_2(x) &= l = w \cdot x + b\end{aligned}$$

Question 4.2

6 P.

Solution 4.2

One point each:

$$\begin{aligned}\frac{\partial L}{\partial m} &= \frac{\partial L}{\partial \hat{y}_1} \frac{\partial \hat{y}_1}{\partial m} = \frac{\partial L}{\partial \hat{y}_1} \sigma'(x + wx + b) \\\frac{\partial L}{\partial l} &= \frac{\partial L}{\partial \hat{y}_2} + \frac{\partial L}{\partial m} \\\frac{\partial L}{\partial b} &= \frac{\partial L}{\partial l} \frac{\partial l}{\partial b} = \frac{\partial L}{\partial l} \cdot 1 \\\frac{\partial L}{\partial w} &= \frac{\partial L}{\partial l} \cdot x \\\frac{\partial L}{\partial k} &= \frac{\partial L}{\partial l} \frac{\partial l}{\partial k} = \frac{\partial L}{\partial l} \cdot w \\\frac{\partial L}{\partial x} &= \frac{\partial L}{\partial k} + \frac{\partial L}{\partial m}\end{aligned}$$

5 Reinforcement learning (9P)

Question 5.1

1 P.

Solution 5.1

action: going from on student to the next. (0.5p)

reward: points for finishing a request. (0.5p)

Question 5.2

2 P.

Solution 5.2

Y maximum points (1p), X greedy algorithm (1p).

0	1	0	0	2
Y	Y	1	0	0
0	Y	0	X	0
1	F	X	X	1
0	0	0	1	0

Question 5.3

3 P.

Solution 5.3

Only exploitation of the next action is used to get maximum reward which does not lead to a maximum reward over multiple steps. (Cannot look ahead, no state transition included) (1P)

Exploitation: use the known good action (0.5 P)

Exploration: try out new actions (0.5 P)


Epsilon greedy (1.P):

$$\pi(a) = \begin{cases} 1 - \epsilon & \text{if } a = \max_a Q_t(a) \\ \epsilon / (n - 1) & \text{else} \end{cases}$$

Question 5.4

3 P.

Solution 5.4



6 Coding: Optimization (7P)

$$v^{(k)} = \mu v^{(k-1)} - \eta \nabla L(w^{(k)}) \quad (1)$$

$$w^{(k+1)} = w^{(k)} - \mu v^{(k-1)} + (1 + \mu) v^{(k)} \quad (2)$$

Question 6.1

7 P.

```
import numpy as np
```

```
class NAG():
```

```
    def __init__(self, learning_rate, momentum=0.9):
        self.learning_rate = learning_rate    (0.5p)
        self.momentum_gradient = None    (0.5p)
        self.momentum = momentum    (0.5p)
        # Initialize momentum gradient with None or zero-like tensor.

    def calculate_update(self, weight_tensor, gradient_tensor):

        if self.momentum_gradient is None:
            self.momentum_gradient = np.zeros_like(weight_tensor)    (1p)
        )

        previous_direction = self.momentum_gradient    (1p)
        self.momentum_gradient = self.momentum * previous_direction -
            learning_rate * gradient_tensor    (1.5p, decide by
            yourself)
        # Compute the current momentum gradient -> 2 P.
        return weight_tensor - self.momentum * previous_direction + (1+
            self.momentum)*self.momentum_gradient    (2p)
```

7 Coding: Pytorch (7P)

Question 7.1

7P.

```
class Inception(nn.Module):
    def __init__(self):
        self.convA2 = nn.Conv2d(64, 64, kernel_size=3, padding=1,
                                  stride=1)
        self.down1 = nn.Conv2d(64, 128, kernel_size=3, padding=1,
                                 stride=2)

        self.convC = nn.Conv2d(256, 128, kernel_size=3, padding=1)

    def forward(self, x):

        convA1 = F.relu(self.convA1(x))

        # TODO
        convA2 = F.relu(self.convA2(convA1))    (0.5p)
        down1 = F.relu(self.down1(convA2))      (0.5p)

        convB = F.relu(self.convB(down1))       (0.5p)
        down2 = F.relu(self.down2(convB))       (0.5p)

        conv_up1 = F.relu(self.conv_up1(down2)) (0.5p)
        conv_all = torch.cat([convB, conv_up1], dim=1) (0.5p)

        convC = F.relu(self.convC(conv_all))    (0.5p)

        conv_up2 = F.relu(self.conv_up2(convC)) (0.5p)
        convD1 = F.relu(self.convD1(conv_up2))  (0.5p)
        convD2 = F.relu(self.convD2(convD1))    (0.5p)
```